


Upload the Dataset

```
from google.colab import files
uploaded = files.upload()
```

 Choose Files churn_prediction (1).csv


- **churn_prediction (1).csv**(text/csv) - 2516120 bytes, last modified: 4/23/2025 - 100% done

Saving churn_prediction (1).csv to churn_prediction (1).csv

Load the Dataset

```
import pandas as pd

# Replace with your actual filename
df = pd.read_csv('/content/churn_prediction (1).csv')
df.head()
```



	customer_id	vintage	age	gender	dependents	occupation	city	customer_nw_category	branch_code	days_since
0	1	3135	66	0	0.0	0	187.0	2	755	
1	6	2531	42	0	2.0	0	1494.0	3	388	
2	7	263	42	1	0.0	0	1096.0	2	1666	
3	8	5922	72	0	0.0	1	1020.0	1	1	
4	9	1145	46	0	0.0	0	623.0	2	317	


5 rows × 21 columns

Data Exploration

```
# Basic info
df.info()

# Descriptive statistics
df.describe()

# Preview column names
print("Columns:", df.columns.tolist())
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22067 entries, 0 to 22066
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   customer_id                          22067 non-null  int64
1   vintage                              22067 non-null  int64
2   age                                  22067 non-null  int64
3   gender                               22067 non-null  int64
4   dependents                           22067 non-null  float64
5   occupation                           22067 non-null  int64
6   city                                 22067 non-null  float64
7   customer_nw_category                 22067 non-null  int64
8   branch_code                          22067 non-null  int64
9   days_since_last_transaction          22067 non-null  float64
10  current_balance                      22067 non-null  float64
11  previous_month_end_balance           22067 non-null  float64
```

```

12 average_monthly_balance_prevQ 22067 non-null float64
13 average_monthly_balance_prevQ2 22067 non-null float64
14 current_month_credit           22067 non-null float64
15 previous_month_credit          22067 non-null float64
16 current_month_debit            22067 non-null float64
17 previous_month_debit           22067 non-null float64
18 current_month_balance          22067 non-null float64
19 previous_month_balance         22067 non-null float64
20 churn                          22067 non-null int64
dtypes: float64(13), int64(8)
memory usage: 3.5 MB
Columns: ['customer_id', 'vintage', 'age', 'gender', 'dependents', 'occupation', 'city', 'customer_nw_category', '

```

Check for Missing Values and Duplicates

```

# Check for missing values
print(df.isnull().sum())

# Check for duplicates
print("Duplicate Rows:", df.duplicated().sum())

```

```

→ customer_id      0
vintage            0
age                0
gender             0
dependents         0
occupation         0
city               0
customer_nw_category 0
branch_code        0
days_since_last_transaction 0
current_balance    0
previous_month_end_balance 0
average_monthly_balance_prevQ 0
average_monthly_balance_prevQ2 0
current_month_credit 0
previous_month_credit 0
current_month_debit 0
previous_month_debit 0
current_month_balance 0
previous_month_balance 0
churn              0
dtype: int64
Duplicate Rows: 0

```

Visualize a Few Features

```

import seaborn as sns
import matplotlib.pyplot as plt

# Set Seaborn style for better visuals
sns.set(style="whitegrid")

# Check if 'Gender' and 'Age' columns exist
if 'Gender' in df.columns:
    plt.figure(figsize=(6, 4))
    sns.countplot(data=df, x='Gender', palette='Set2')
    plt.title('Gender Distribution')
    plt.xlabel('Gender')
    plt.ylabel('Count')
    plt.show()
else:
    print("Column 'Gender' not found in DataFrame.")

```

```

if 'Age' in df.columns:
    plt.figure(figsize=(6, 4))
    sns.histplot(df['Age'], kde=True, color='skyblue', bins=30)
    plt.title('Age Distribution')
    plt.xlabel('Age')
    plt.ylabel('Frequency')
    plt.show()
else:
    print("Column 'Age' not found in DataFrame.")

```

Column 'Gender' not found in DataFrame.
Column 'Age' not found in DataFrame.

Identify Target and Features

```
print(df.columns.tolist())
```

```

target_column = 'churn'
X = df.drop(target_column, axis=1)
y = df[target_column]
df.head()

```

['customer_id', 'vintage', 'age', 'gender', 'dependents', 'occupation', 'city', 'customer_nw_category', 'branch_code', 'days_since']

	customer_id	vintage	age	gender	dependents	occupation	city	customer_nw_category	branch_code	days_since
0	1	3135	66	0	0.0	0	187.0	2	755	
1	6	2531	42	0	2.0	0	1494.0	3	388	
2	7	263	42	1	0.0	0	1096.0	2	1666	
3	8	5922	72	0	0.0	1	1020.0	1	1	
4	9	1145	46	0	0.0	0	623.0	2	317	

5 rows × 21 columns

Convert Categorical Columns to Numerical

```

# Identify categorical columns
cat_cols = X.select_dtypes(include='object').columns
print("Categorical Columns:", cat_cols.tolist())

# Apply label encoding temporarily (can be replaced with OneHot later)
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()

for col in cat_cols:
    X[col] = le.fit_transform(X[col])

```

Categorical Columns: []

One-Hot Encoding

```
X = pd.get_dummies(X, drop_first=True)
```

Feature Scaling

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Train-Test Split

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)
```

Model Building

```
from sklearn.ensemble import RandomForestClassifier

model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
```




Evaluation

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
plt.show()
```

Accuracy: 0.8568192115994563

	precision	recall	f1-score	support
0	0.87	0.96	0.91	3544
1	0.73	0.44	0.55	870
accuracy			0.86	4414
macro avg	0.80	0.70	0.73	4414
weighted avg	0.85	0.86	0.84	4414



Make Prediction

```
print(X.columns.tolist())

import pandas as pd

# Create a dictionary that includes all columns from training
new_input_dict = {
    'Age': 35,
    'Gender_Male': 1,
    'Gender_Female': 0,
    'Plan_Basic': 0,
    'Plan_Premium': 1,
    'MonthlyCharges': 5000,
    # ... include all other one-hot encoded or numeric features, set missing to 0
}

# Convert to DataFrame
new_input_df = pd.DataFrame([new_input_dict])

# Reindex to match training column order
new_input_df = new_input_df.reindex(columns=X.columns, fill_value=0)

# Scale
new_input_scaled = scaler.transform(new_input_df)

# Predict
prediction = model.predict(new_input_scaled)
print("Prediction:", "Churn" if prediction[0] == 1 else "Not Churn")
```

```

[ 'customer_id', 'vintage', 'age', 'gender', 'dependents', 'occupation', 'city', 'customer_nw_category', 'branch_co
Prediction: Churn

```

Convert to DataFrame and Encode

```

input_dict = {
    'Age': [35],
    'Gender': ['Male'],
    'Plan': ['Basic'],
    # Add more fields as per your original dataset
}

input_df = pd.DataFrame(input_dict)

# Convert categorical variables
for col in input_df.select_dtypes(include='object'):
    input_df[col] = le.fit_transform(input_df[col])

# Align columns
input_df = pd.get_dummies(input_df)
input_df = input_df.reindex(columns=X.columns, fill_value=0)

input_scaled = scaler.transform(input_df)

```

Predict the Final Grade

```

final_prediction = model.predict(input_scaled)
print("Final Prediction:", final_prediction)

```

```

Final Prediction: [1]

```

Deployment - Building an Interactive App

```

# Simulate form input in Colab
user_input = pd.DataFrame({
    'Age': [30],
    'Gender': ['Male'],
    'Plan': ['Premium'],
    # Add other features...
})

# Encode, align, scale
for col in user_input.select_dtypes(include='object'):
    user_input[col] = le.fit_transform(user_input[col])

user_input = pd.get_dummies(user_input)
user_input = user_input.reindex(columns=X.columns, fill_value=0)

user_input_scaled = scaler.transform(user_input)
prediction = model.predict(user_input_scaled)

print("Prediction:", "Churn" if prediction[0] == 1 else "Not Churn")

```

```

Prediction: Churn

```

Create a Prediction Function

```
def preprocess_input(input_data, scaler, encoder, base_columns):
    """
    Preprocess input data: encode, one-hot, scale, and align columns.

    Args:
    - input_data (pd.DataFrame): Raw input data.
    - scaler (StandardScaler): Fitted scaler.
    - encoder (LabelEncoder): Fitted label encoder for categorical vars.
    - base_columns (list): List of original X.columns after one-hot.

    Returns:
    - np.array: Scaled and aligned feature vector.
    """
    data = input_data.copy()

    for col in data.select_dtypes(include='object').columns:
        data[col] = encoder.fit_transform(data[col])

    data = pd.get_dummies(data)
    data = data.reindex(columns=base_columns, fill_value=0)
    data_scaled = scaler.transform(data)

    return data_scaled

def predict_churn(input_dict, model, scaler, encoder, base_columns):
    """
    Make churn prediction from raw input dictionary.

    Args:
    - input_dict (dict): User inputs as key-value pairs.
    - model (trained model): Trained classifier.
    - scaler (StandardScaler): Trained scaler.
    - encoder (LabelEncoder): Trained label encoder.
    - base_columns (list): Reference for column alignment.

    Returns:
    - str: Prediction result.
    """
    input_df = pd.DataFrame([input_dict])
    processed = preprocess_input(input_df, scaler, encoder, base_columns)
    prediction = model.predict(processed)[0]
    return "Churn" if prediction == 1 else "Not Churn"
```

create the gradio interface

```
!pip install -q gradio
import gradio as gr

def predict_churn(age, gender, plan, monthly_charges):
    # Create input DataFrame
    input_dict = {
        'Age': [age],
        'Gender': [gender],
        'Plan': [plan],
        'MonthlyCharges': [monthly_charges]
    }
    input_df = pd.DataFrame(input_dict)

    # Encode
```

```

for col in input_df.select_dtypes(include='object'):
    input_df[col] = le.fit_transform(input_df[col])

# One-hot encoding (if needed)
input_df = pd.get_dummies(input_df)
input_df = input_df.reindex(columns=column_names, fill_value=0)

# Scale
input_scaled = scaler.transform(input_df)

# Predict
prediction = model.predict(input_scaled)[0]
return "Churn" if prediction == 1 else "Not Churn"

```




54.1/54.1 MB	17.5 MB/s	eta 0:00:00
322.9/322.9 kB	25.9 MB/s	eta 0:00:00
95.2/95.2 kB	8.2 MB/s	eta 0:00:00
11.5/11.5 MB	112.1 MB/s	eta 0:00:00
72.0/72.0 kB	6.2 MB/s	eta 0:00:00
62.5/62.5 kB	5.6 MB/s	eta 0:00:00

```

iface = gr.Interface(
    fn=predict_churn,
    inputs=[
        gr.Number(label="Age"),
        gr.Dropdown(choices=["Male", "Female"], label="Gender"),
        gr.Dropdown(choices=["Basic", "Premium", "Gold"], label="Plan"),
        gr.Number(label="Monthly Charges")
    ],
    outputs="text",
    title="Churn Prediction App",
    description="Enter details to predict if a customer will churn."
)

iface.launch()

```


 It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sharing must be e